



Inspiring Excellence

Network Layer

Routing Algorithm

Link State Routing

Lecture 12 | CSE421 – Computer Networks

Department of Computer Science and Engineering

School of Data & Science

Chapter 4: Network Layer

- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- 4.5 Routing algorithms
 - Distance Vector
 - Link state
 - Hierarchical routing
- 4.6 Routing in the Internet
 - RIP
 - OSPF
 - BGP

Link-State Routing Protocols



- **Distance Vector** routing protocols are like road signs.
 - Routers must make preferred path decisions based on a distance or metric to a network.
- **Link-State** routing protocols are more like a road map.
 - They create a topological map of the network and each router uses this map to determine the shortest path to each network.


Link-State Routing Protocols

- Centralized Routing Algorithm
 - computes the least-cost path using complete, global knowledge about the network.
- Link-state routing protocols are also known as **shortest path first protocols**
 - The Link state routing protocol uses Dijkstra's algorithm which is used to find the shortest path from one node to every other node in the network.
- While they have the reputation of being much more complex than distance vector, the basic functionality and configuration of link state routing protocols are not complex.

A Link-State Routing Algorithm


- **Centralized Routing Algorithm**
- **Uses Dijkstra's algorithm- Shortest Path First**
 - net topology, link costs known to all nodes
 - computes least cost paths from one node ("source") to all other nodes
 - gives **forwarding/routing table** for that node
 - iterative: after k iterations, know least cost path to k destinations

SPF Algorithm - Example




- Dijkstra's algorithm is commonly referred to as the Shortest Path First (**SPF**) algorithm.
- This algorithm **accumulates costs** along each path, from source to destination.

SPF Algorithm



- To illustrate how SPF operates, each path in the figure is labeled with an arbitrary value for **cost**.
- Each router calculates the SPF algorithm and determines the cost of a link **from its own perspective**.


SPF Algorithm



- For example:
 - The cost of the shortest path for R2 to send packets to the LAN attached to R3 is 27 ($20 + 5 + 2 = 27$).

SPF Algorithm

R1 uses 3 “hops”
but “faster” links.



- R1 has data to send to the network on R5.
 - You might think that R1 would send directly to R4 (2 hops) instead of to R3 (3 hops).


Link-State Routing Process

5 Step Process

1. **Each router** learns about its own **directly connected networks**.
2. **Each router** is responsible for **contacting its neighbors (exchange Hello packet)** on directly connected networks.
3. **Each router** builds a **link-state packet (LSP)** containing the state of each directly connected link.
4. **Each router** **floods the LSP to all routers**, who then store all LSPs received in a database.
5. **Each router** uses the LSPs to **construct a database** that is a **complete map of the topology and computes the best path** to each destination network.


Step 1: Directly Connected Networks

We will focus on the Link-State routing process from the perspective of R1.



- Each router learns about its own directly connected networks.
- When a router interface is configured with an IP address and subnet mask and activated, the interface becomes part of that network.
- Regardless of the routing protocols used, these directly connected networks are now part of the routing table.


Step 2: Hello Packets



- Each router is responsible for contacting its neighbors on directly connected networks.
- The router will not be aware of any neighbor routers on the link until it receives a Hello packet from that neighbor.

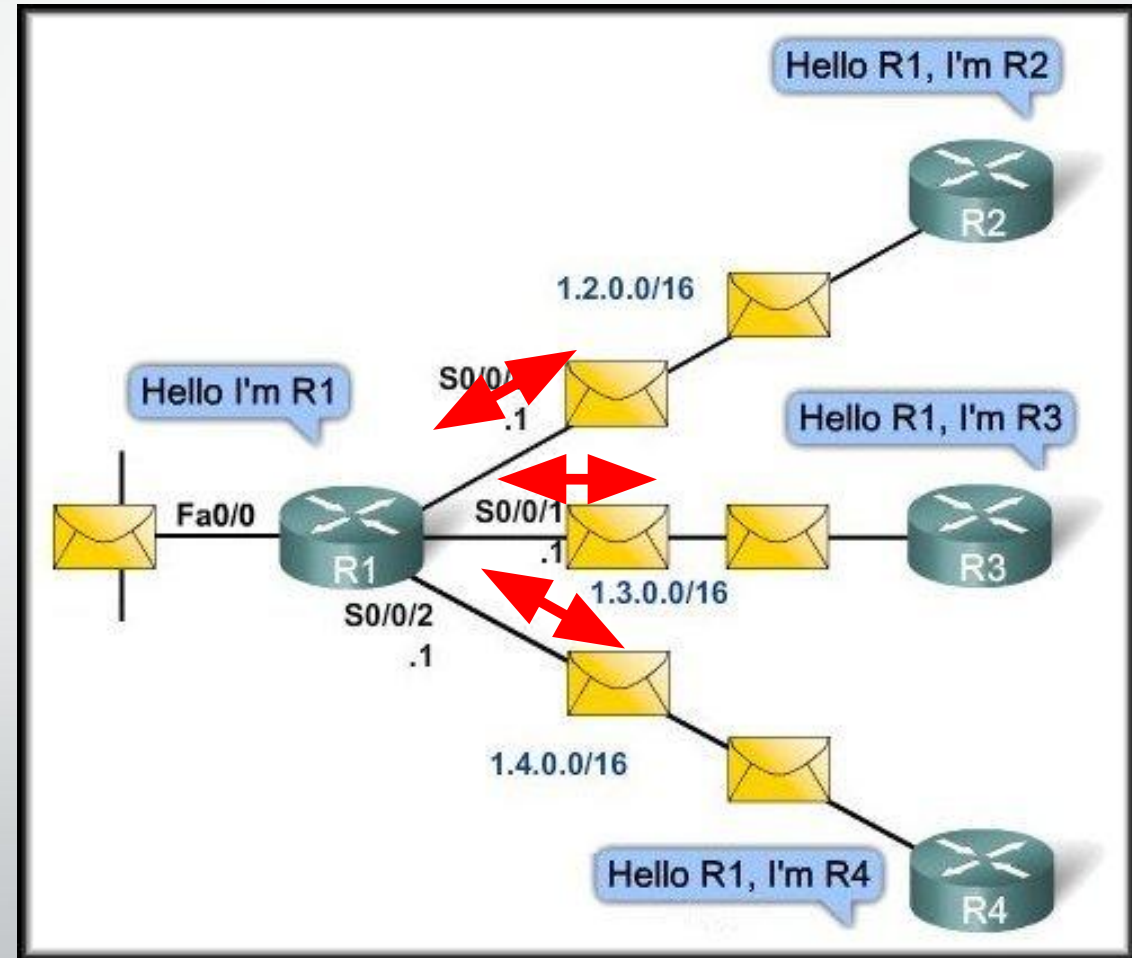
At that time, it establishes an adjacency with the neighboring router.

Step 2: Hello Packets

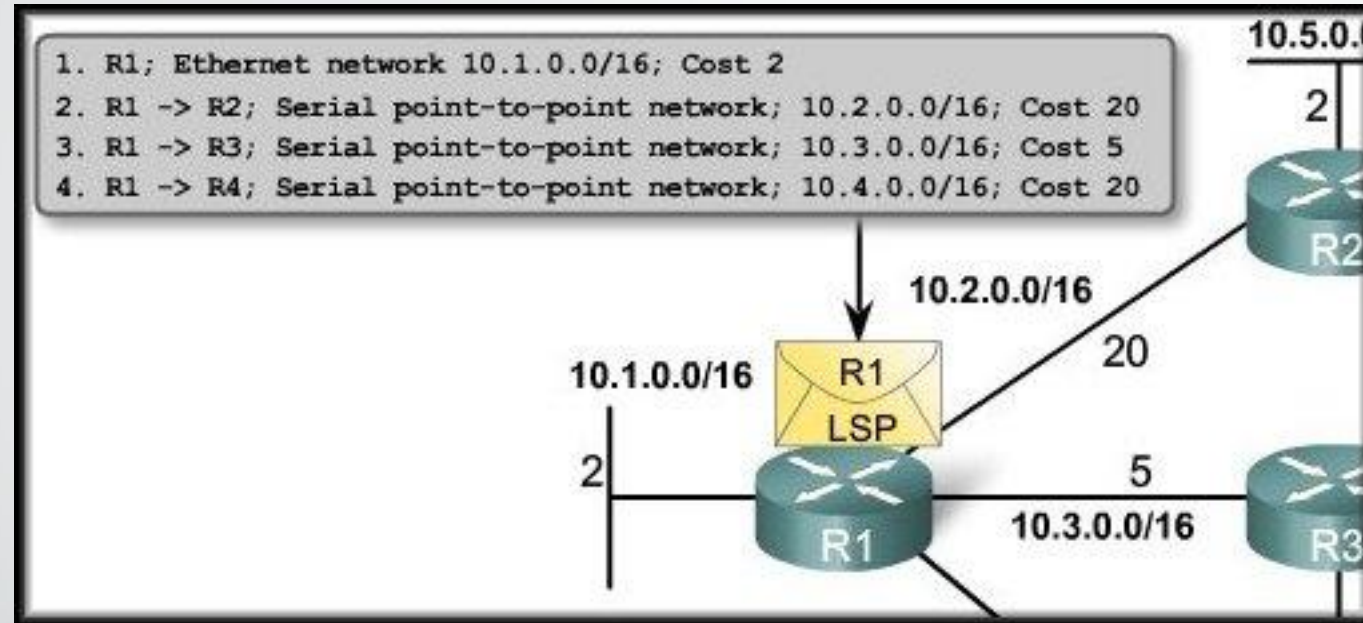


Step 2: Hello Packets

- A **neighbor** is any other router that is enabled with the **same link-state routing protocol**.
- These small Hello packets continue to be exchanged between two adjacent neighbors.
- These packets serve as a **keep alive** function to monitor the state of the neighbor.



Step 3: Build the Link-State Packet




- Each router builds a link-state packet (LSP) containing the state of each directly connected link.
- The LSP contains the link-state information about the sending router's links.
- The router only sends LSPs out interfaces where it has established adjacencies with other routers.

Step 4: Flooding Link-State Packets


No response to Hello
– no LSP sent.

Flooding of R1 LSP




- Each router floods the LSP to all neighbors, who then store all LSPs received in a database.
- Whenever a router receives an LSP from a neighboring router, it immediately sends that LSP out all other interfaces, **except the interface that received the LSP.**

Step 4: Flooding Link-State Packets




- Link-state routing protocols calculate the SPF algorithm **after the flooding** is complete.
- As a result, link-state routing protocols **reach convergence much faster** than distance vector routing protocols.

Step 4: Flooding Link-State Packets



- An **LSP** needs to be sent only:
 - During **initial startup** of the router or routing protocol.
 - Whenever there is a **change in the topology** (link going down or coming up) or a neighbor adjacency being established or broken.

Step 5: Constructing a Link-State Database



- Each router uses the LSPs to **construct a database** that is a **complete map of the topology** and computes the **best path** to each destination network.

R1: Building the SPF Tree

R1 Link State Database

R1 Links-states:

- Connected to neighbor R2 on network 10.2.0.0/16, cost of 20
- Connected to neighbor R3 on network 10.3.0.0/16, cost of 5
- Connected to neighbor R4 on network 10.4.0.0/16, cost of 20
- Has a network 10.1.0.0/16, cost of 2

LSPs from R2:

- Connected to neighbor R1 on network 10.2.0.0/16, cost of 20
- Connected to neighbor R5 on network 10.9.0.0/16, cost of 10
- Has a network 10.5.0.0/16, cost of 2

LSPs from R3:


- Connected to neighbor R1 on network 10.3.0.0/16, cost of 5
- Connected to neighbor R4 on network 10.7.0.0/16, cost of 10
- Has a network 10.6.0.0/16, cost of 2

LSPs from R4:

- Connected to neighbor R1 on network 10.4.0.0/16, cost of 20
- Connected to neighbor R3 on network 10.7.0.0/16, cost of 10
- Connected to neighbor R5 on network 10.10.0.0/16, cost of 10
- Has a network 10.8.0.0/16, cost of 2

LSPs from R5:

- Connected to neighbor R2 on network 10.9.0.0/16, cost of 10
- Connected to neighbor R4 on network 10.10.0.0/16, cost of 10
- Has a network 10.11.0.0/16, cost of 2



Step 5: Constructing a Link-State Database

R1's Link-State Database

LSPs from R2:

- Connected to neighbor R1 on network 10.2.0.0/16, cost of 20
- Connected to neighbor R5 on network 10.9.0.0/16, cost of 10
- Has a network 10.5.0.0/16, cost of 2

LSPs from R3:

- Connected to neighbor R1 on network 10.3.0.0/16, cost of 5
- Connected to neighbor R4 on network 10.7.0.0/16, cost of 10
- Has a network 10.6.0.0/16, cost of 2

LSPs from R4:


- Connected to neighbor R1 on network 10.4.0.0/16, cost of 20
- Connected to neighbor R3 on network 10.7.0.0/16, cost of 10
- Connected to neighbor R5 on network 10.10.0.0/16, cost of 10
- Has a network 10.8.0.0/16, cost of 2

LSPs from R5:

- Connected to neighbor R2 on network 10.9.0.0/16, cost of 10
- Connected to neighbor R4 on network 10.10.0.0/16, cost of 10
- Has a network 10.11.0.0/16, cost of 2

R1 Link-states:

- Connected to neighbor R2 on network 10.2.0.0/16, cost of 20
- Connected to neighbor R3 on network 10.3.0.0/16, cost of 5
- Connected to neighbor R4 on network 10.4.0.0/16, cost of 20
- Has a network 10.1.0.0/16, cost of 2




- As a result of the flooding process, router R1 has learned the link-state information for each router in its routing area.

Step 5: Constructing a Link-State Database


| Destination | Shortest Path | Cost |
|-------------|----------------------|------|
| R2 LAN | R1 to R2 | 22 |
| R3 LAN | R1 to R3 | 7 |
| R4 LAN | R1 to R3 to R4 | 17 |
| R5 LAN | R1 to R3 to R4 to R5 | 27 |

Each router in the topology determines the shortest path **from its own perspective**.



- With a complete link-state database, R1 can now use the database and the shortest path first (SPF) algorithm to calculate the preferred path or **shortest path** to each network.

R1: Building the SPF Tree




- All LSPs have been processed using the SPF algorithm and R1 has now constructed the complete SPF tree.

Generating a Routing Table

SPF Information

- Network 10.5.0.0/16 via R2 serial 0/0/0 at a cost of 22
- Network 10.6.0.0/16 via R3 serial 0/0/1 at a cost of 7
- Network 10.7.0.0/16 via R3 serial 0/0/1 at a cost of 15
- Network 10.8.0.0/16 via R3 serial 0/0/1 at a cost of 17
- Network 10.9.0.0/16 via R2 serial 0/0/0 at a cost of 30
- Network 10.10.0.0/16 via R3 serial 0/0/1 at a cost of 25
- Network 10.11.0.0/16 via R3 serial 0/0/1 at a cost of 27



Building Routing Table

- Creation of the states of the links by each node, called the link state packets (LSP)
- Dissemination of LSPs to every other routers, called flooding (efficiently)
- Constructing a Link-State Database using LSPs
- Formation of a shortest path tree for each node
- Calculation of a routing table based on the shortest path tree

Advantages: Link-State Routing

- **Fast Network Convergence**– On receiving an LSP, link-state routing protocols immediately flood the LSP out all interfaces without any changes except for the interface from which the LSP was received.
- **Topological Map**– Using the SPF tree, each router can separately determine the shortest path to every network.
- **Hierarchical Design**– Link-state routing protocols use multiple areas and create a hierarchical design to networks areas. The multiple areas allow better route summarization.
- **Event-driven Updates**– After initial flooding of LSPs, the LSPs are sent only when there is a change in the topology and contain only the information regarding that change. The LSP contains only the information about the affected link. The link-state never sends periodic updates.

Comparison

| | Distance Vector | Link State |
|-------------------------------|--|---|
| Network view | Topology knowledge from the neighbor point of view | Common and complete knowledge of the network topology |
| Best Path Calculation | Based on fewest number of hops | Based on the link cost |
| Updates | Full routing table | Link State Updates |
| Algorithm | Bellman-Ford | Dijkstra |
| CPU and Memory | Low utilization | Intensive |
| hierarchical Structure | No | Yes |
| Convergence time | Moderate | Fast |



THE END