



Transport Layer (TCP)

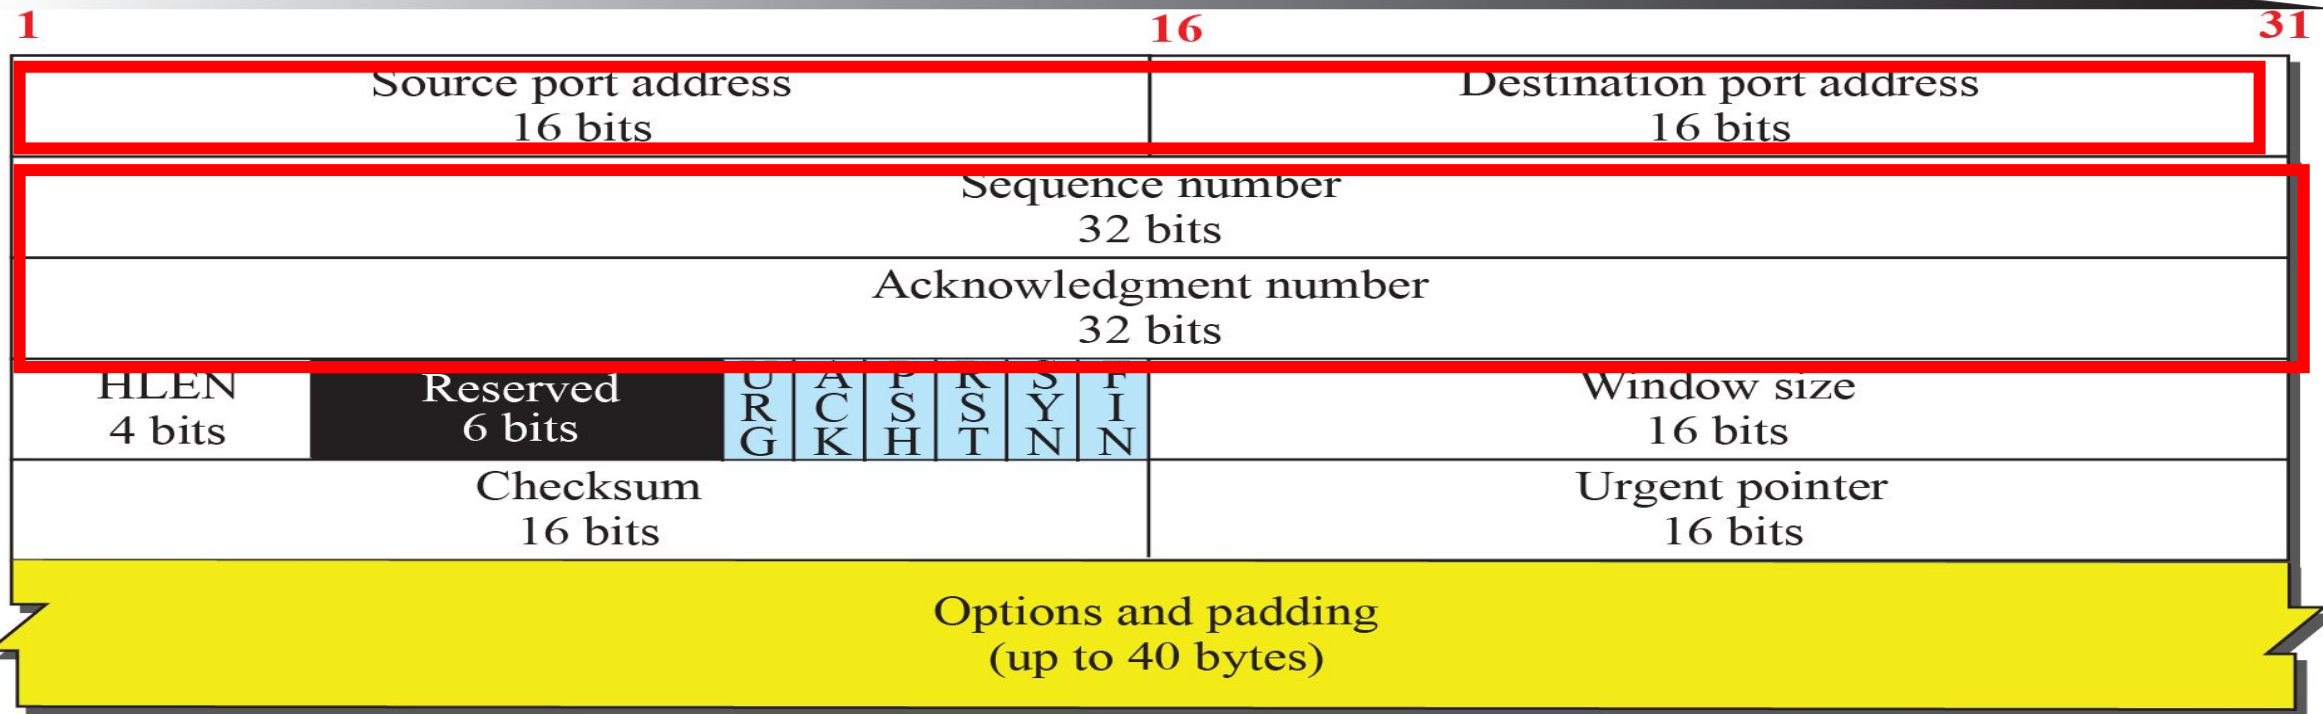
Lecture 5 | CSE421 – Computer Networks

Department of Computer Science and Engineering
School of Data & Science

Objectives

- TCP Header
- TCP Services

TCP Segment Header



b. Header

Byte Number

- The bytes of data being transferred in each connection are numbered by TCP.
- The numbering starts with an arbitrarily generated number.
- An arbitrary number between **0 and $2^{32} - 1$** for the number of the first byte.
- For example if the number of the first byte happens to be **1067** and the total data to be sent is **3000 bytes**.
- What is the byte number for the first byte of data and last byte of data?

First Byte Number
1067



Last Byte Number
4066

Sequence Numbers

- The sequence number of the first segment is **the ISN (initial sequence number)**, which is a random number (byte number).
- The sequence number of any other segment is the sequence number of the previous segment plus the number of bytes (**real or imaginary**) carried by the previous segment.
- Suppose a TCP connection is transferring a file of 5,000 bytes. The first byte is numbered 10,001. **What are the sequence numbers for each segment if data are sent in five segments, each carrying 1,000 bytes?**

- **Solution:**

Segment 1	→	Sequence Number:	10001	Range:	10001	to	11000
Segment 2	→	Sequence Number:	11001	Range:	11001	to	12000
Segment 3	→	Sequence Number:	12001	Range:	12001	to	13000
Segment 4	→	Sequence Number:	13001	Range:	13001	to	14000
Segment 5	→	Sequence Number:	14001	Range:	14001	to	15000


Acknowledgement Number

- If receiving host TCP receives uncorrupted data, then...
- It is acknowledged using the acknowledgement number
- The value of the acknowledgment field in a segment defines the **number of the next byte** the receiver expects to receive.
- For example if the sender receives **1001** as the acknowledgement number.
- What does it mean?

Received all data up to 1000, tells the sender that it ready to receive the next data from 1001 byte number.

Note :This does not indicate receiver has received 1000 bytes of data.


Acknowledgement Number



Acknowledgement Number

- The acknowledgment number is **cumulative**.
- Receiver acknowledges multiple data segments in one acknowledgement.


Sender



Receiver




Header Length



- Header Length :
 - Indicates the number of 4-byte words
 - The length of the header can be between 20 and 60 bytes

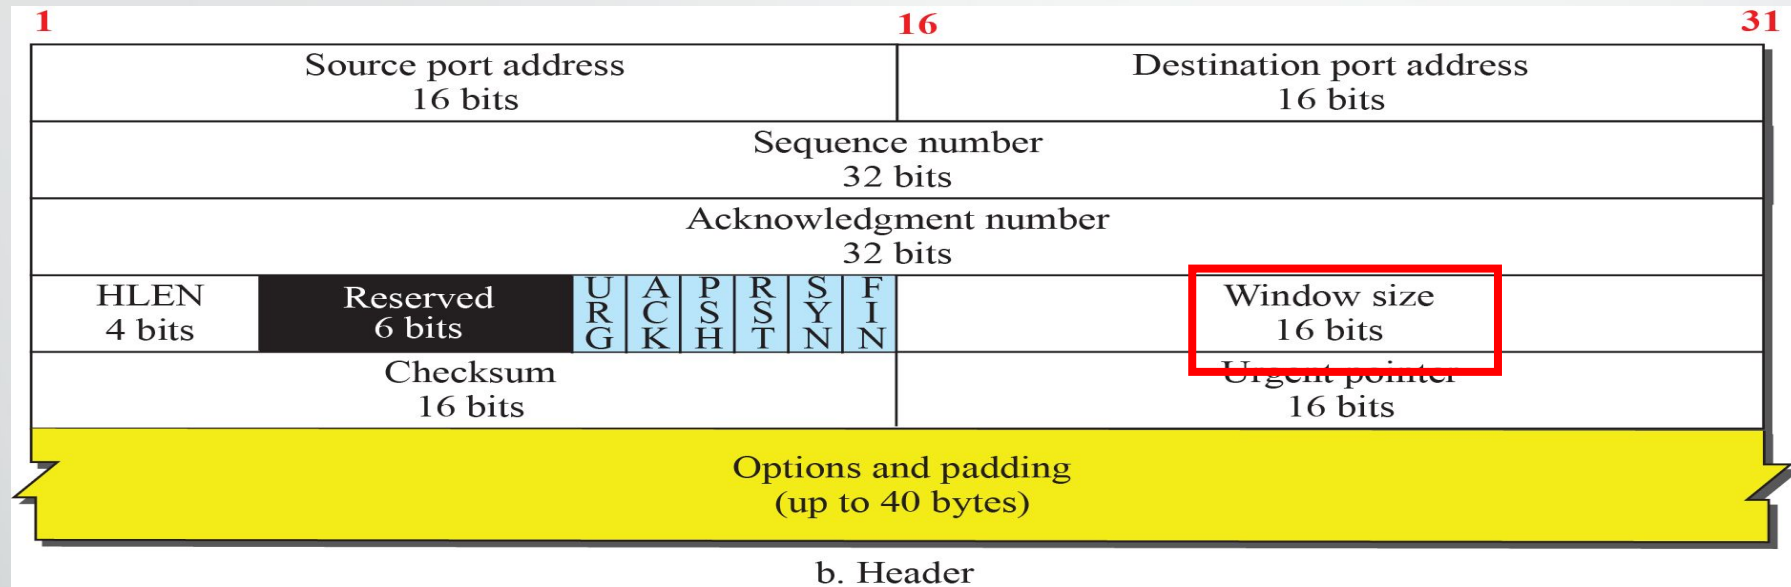
Control Bits



URG: Urgent pointer is valid
ACK: Acknowledgment is valid
PSH : Request for push
RST : Reset the connection
SYN: Synchronize sequence numbers
FIN : Terminate the connection

- Control Bits:
 - This field defines 6 different control bits or flags
 - One or more of these bits can be set at a time
 - These bits help indicate connection establishment and termination, flow control

Window Size



- Window Size:
 - This field defines size of data in bytes of the sending TCP process
 - The maximum size of the window is 65,535 bytes
 - Normally referred to as the receiving window (rwnd)
 - The sender must obey the dictation of the receiver in this case

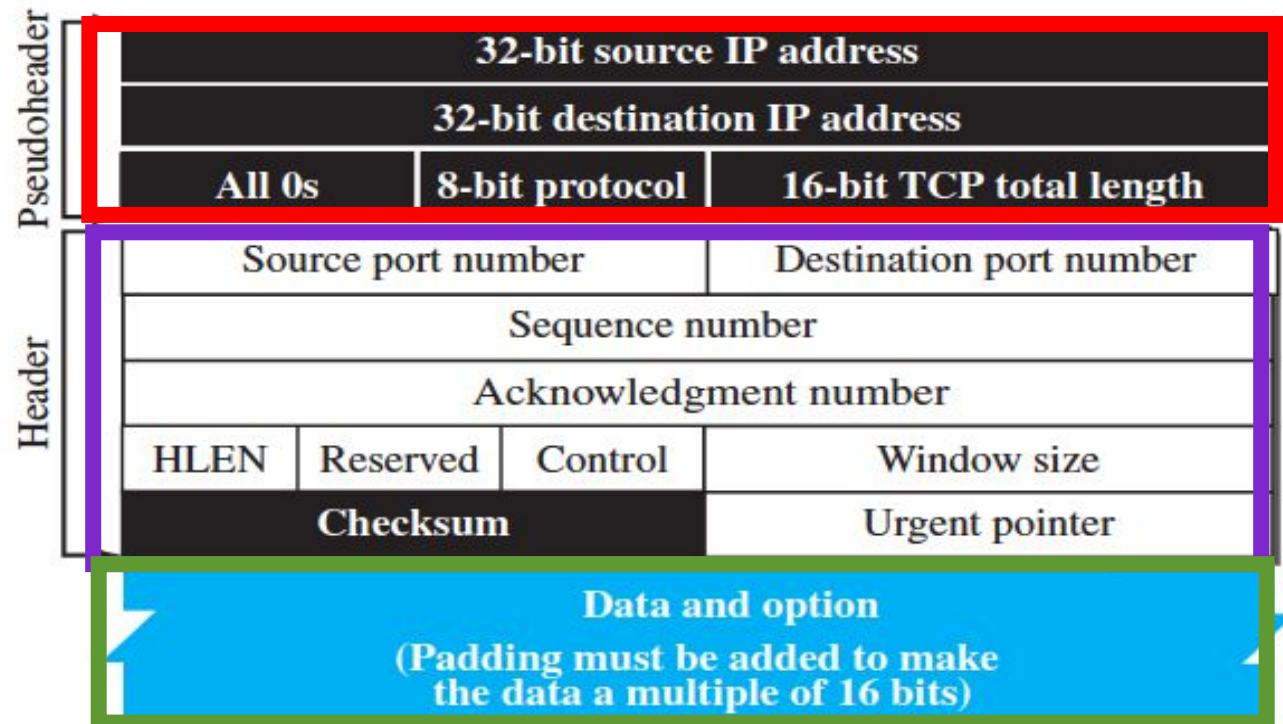
Checksum

- This 16 bits field is used to detect errors (*i.e.*, flipped bits) in the transmitted segment (intentionally or unintentionally) while traveling through the network.
- Also present in UDP header
- Mandatory in TCP but not in UDP
- Process is same for both protocols


UDP HEADER

Source port number 16 bits	Destination port number 16 bits
Total length 16 bits	Checksum 16 bits

- TCP/UDP Header
- TCP/UDP Body
- Pseudo IP Header



Checksum



Checksum Process

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - Not equal - error detected
 - Equal - no error detected. *But maybe errors nonetheless? More later*

Checksum: an example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0	
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1	
<hr/>																	
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
<hr/>																	
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0	
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1	

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Checksum: weak protection!



example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1


Even though numbers have changed (bit flips), *no* change in checksum!

Urgent Pointer

- This 16-bit field, which is valid only if the urgent flag is set.
- Used when the segment contains urgent data.
- It defines a value that must be added to the sequence number to obtain the number of the last urgent byte in the data section of the segment.



Options



- There can be up to 40 bytes of optional information in the TCP header
- Provides a way to deal with limitations of the original header
- For example :
 - MSS (Maximum Segment Size) is defined as the largest block of data that a sender using TCP will send to the receiver

Functions of the Transport Layer

- Primary responsibilities:

1. Segmenting the data and managing each piece.

2. Reassembling the segments into streams of application data.

UDP & TCP

3. Identifying the different applications.


4. Multiplexing

Only TCP

5. Establishing and terminating a connection.

6. Enabling error control and recovery.

7. Performing flow control between end users.

- 
- **Function 6**
Connection Establishment and
Termination for Reliability


Connection Establishment

- TCP sets up a connection between end hosts before sending data
- This process is known as **"Three-way handshake"**
- After the connection is established the hosts can send data

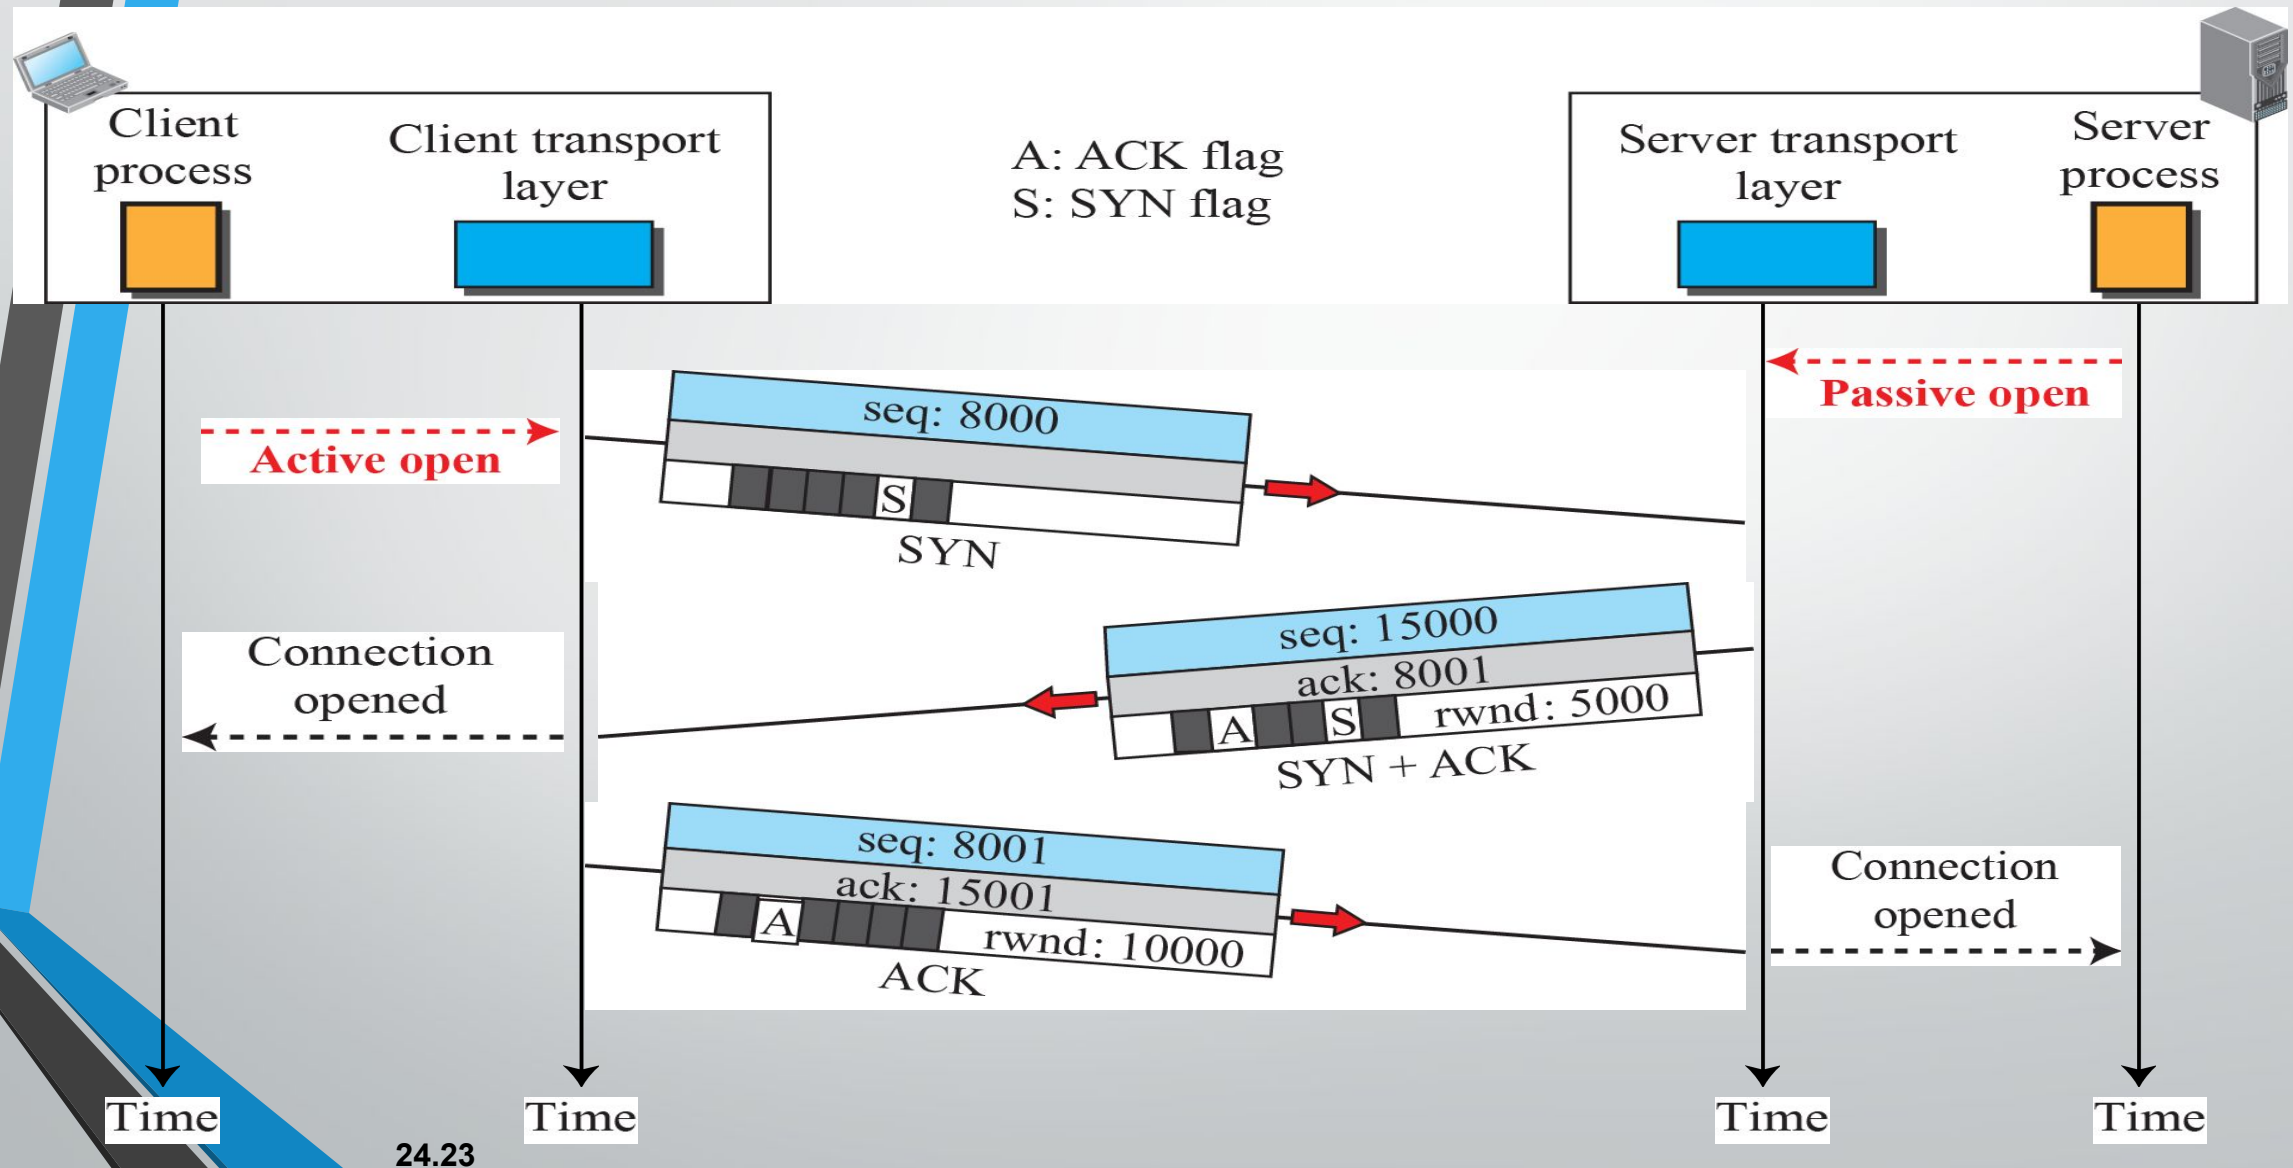
3 Way Handshake : Connection Establishment

Source Port No.	Destination Port No.
Sequence No.	
Acknowledgement No.	
U	A P R S F Window Size
Others	


49152	80	49152
100	0	101
0 1011		
A	S	4000 Bytes 8000 Bytes
Others		



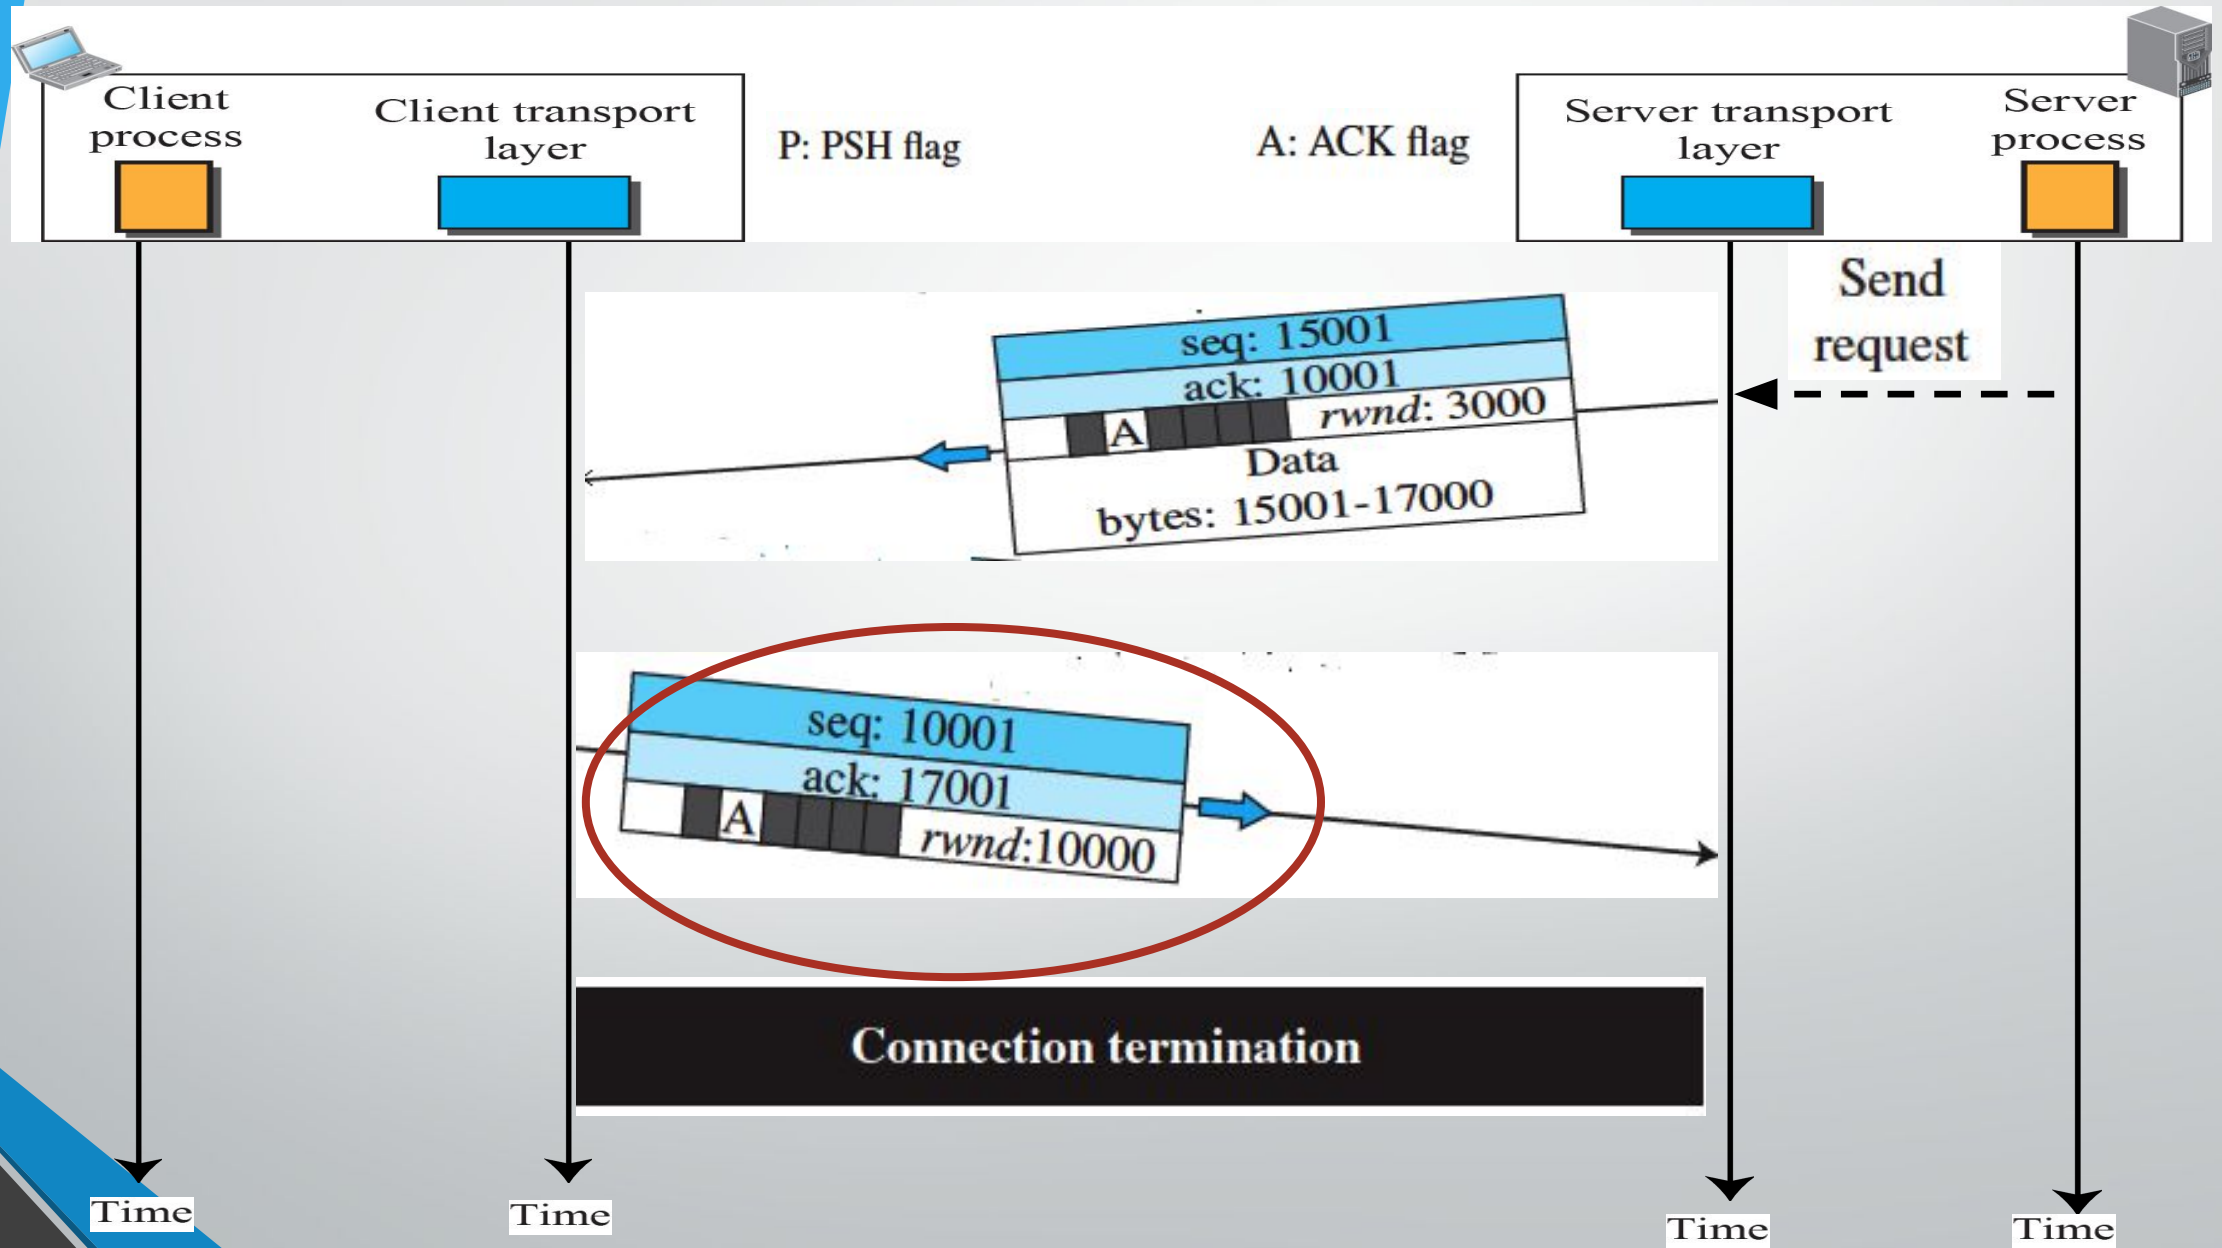
3 Way Handshake : Connection Establishment



Data Transfer




Data Transfer Continued..




3 Way Handshake : Connection Termination

Source Port No.					Destination Port No.				
Sequence No.									
Acknowledgement No.									
	U	A	P	R	S	F	Window Size		
Others									

49152 80					80 49152				
10001 17001					10002				
17001 17002					10002				
		A				F	4000 Bytes 8000 Bytes		
Others									



Connection Termination :: Half Close



- 
- **Function 6**
Error Control and Recovery for Reliability

Reliability in TCP

- TCP provides **reliability** using **error control**
- Error control includes mechanisms for
 - detecting and resending corrupted segments
 - resending lost segments
 - storing out-of order segments until missing segments arrive
 - detecting and discarding duplicated segments.
- Error control in TCP is achieved through
 - **Checksum**
 - **Acknowledgment**
 - **Time-out and retransmission**

Error Control

- Checksum

- Each segment includes a checksum field, which is used to check for a corrupted segment
- If a segment is corrupted, as detected by an invalid checksum, the segment is discarded

- Acknowledgment

- Using Acknowledgement Number to confirm the receipt of data segments.
- To confirm control segments that carry no data, but consume a sequence number
- ACK segments **do not consume sequence numbers** and **are not acknowledged**.

Error Control

- **Retransmission**

- When a segment is sent, it is stored in a queue until it is acknowledged.
- Retransmission of segment will occur

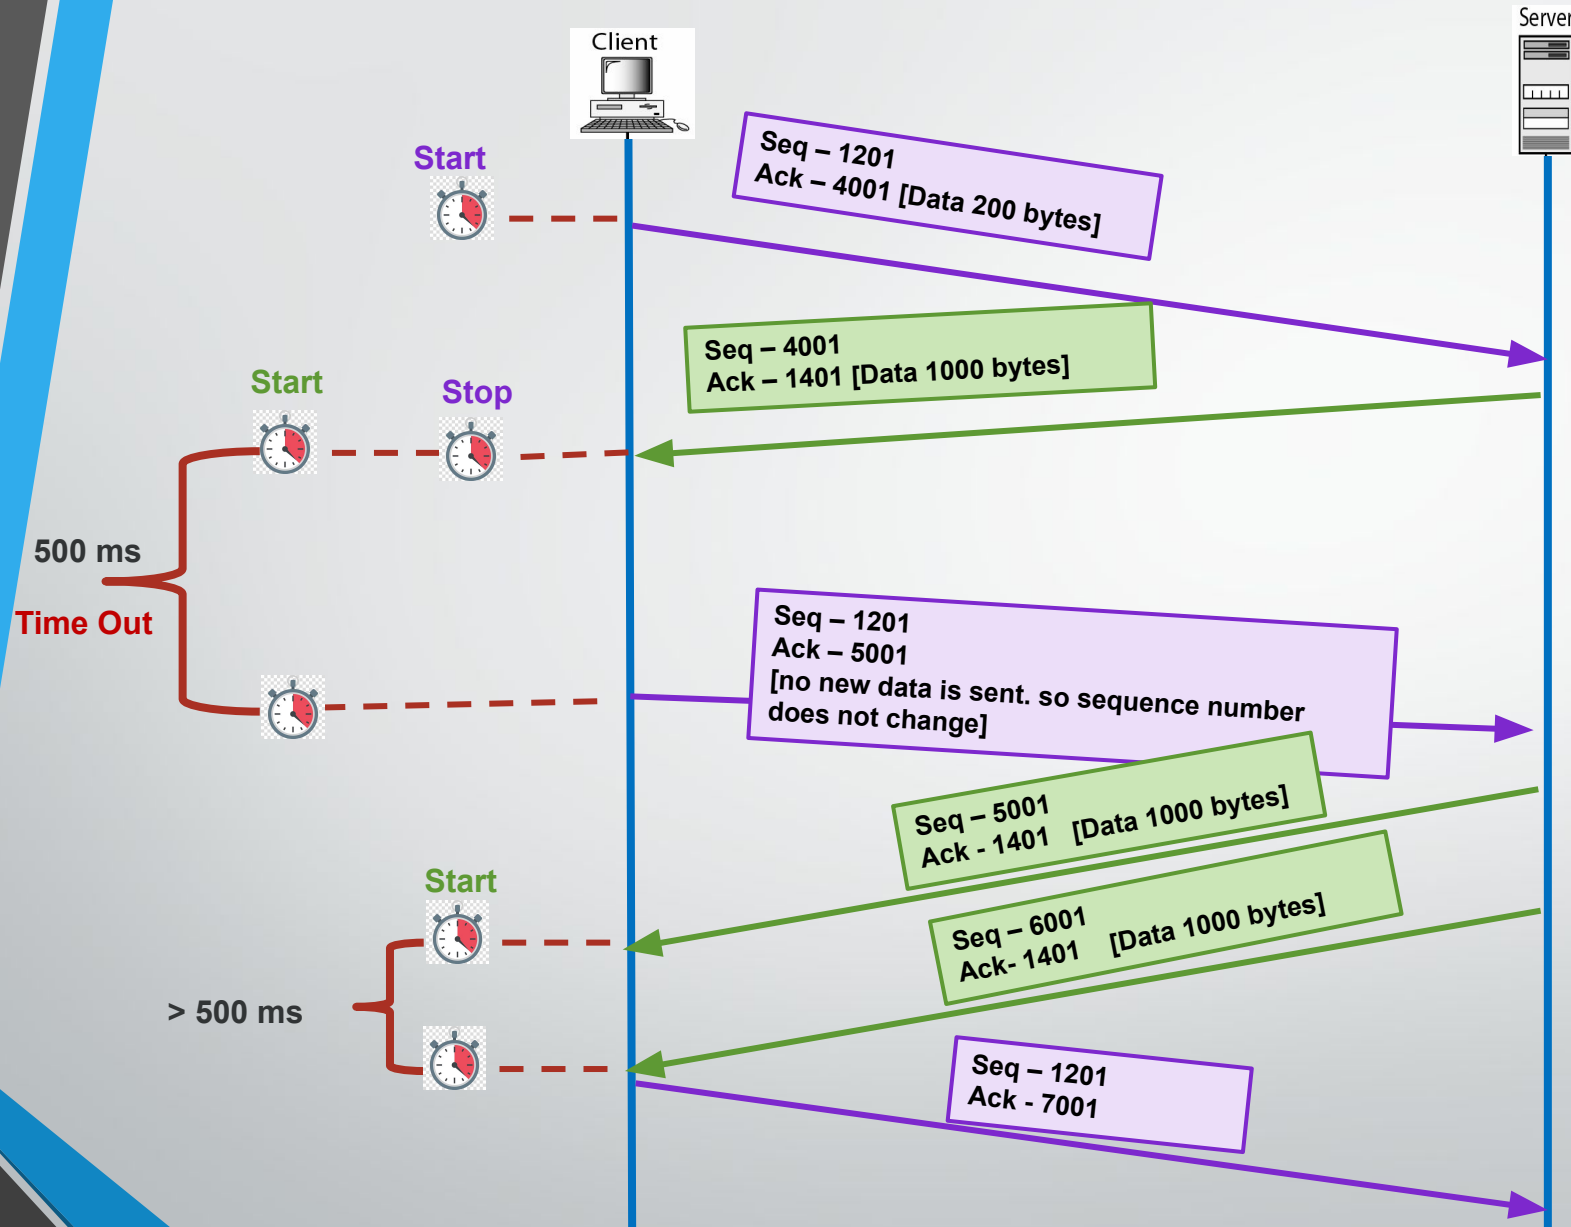
- **After Retransmission Time Out**

- The sending TCP maintains one retransmission time-out (RTO) timer for each connection.
- When the timer matures TCP resends the segment in the front of the queue if the segment is not acknowledged

- **After Three Duplicate ACK Segments**


- To be explained later

Normal Operation



Other Scenarios

- Segment Lost or Corrupted?
- Retransmission of segment ??
- How will the sender know ??
- What about the receiver, not aware of a packet sent?
- **RTO - Retransmission** after time out.




Out of Order Segments

- TCP implementations today do not discard out-of-order segments.
- They store them temporarily .
- Flag them as out-of-order segments until the missing segments arrive.
- Out-of-order segments are never delivered to the process.
- TCP guarantees that data are delivered to the process in order.

Lost ACK

Buffer



Buffer



Seq - 11001
Ack - 1401 [Data 1000 bytes]

Seq - 1401
Ack - 12001

Seq - 11001
Ack - 1401 [Data 1000 bytes]

Seq - 1401
Ack - 12001

Seq - 12001
Ack - 1401 [Data 1000 bytes]

Seq - 13001
Ack - 1401 [Data 1000 bytes]

Seq - 14001
Ack - 1401 [Data 1000 bytes]

Seq - 15001
Ack - 1401 [Data 1000 bytes]

Seq - 12001
Ack - 1401 [Data 1000 bytes]

Seq - 1401
Ack - 12001

Seq - 1401
Ack - 12001

Seq - 1401
Ack - 12001

Seq - 1401
Ack - 12001



Time out & Restart Timer



Time Out

1

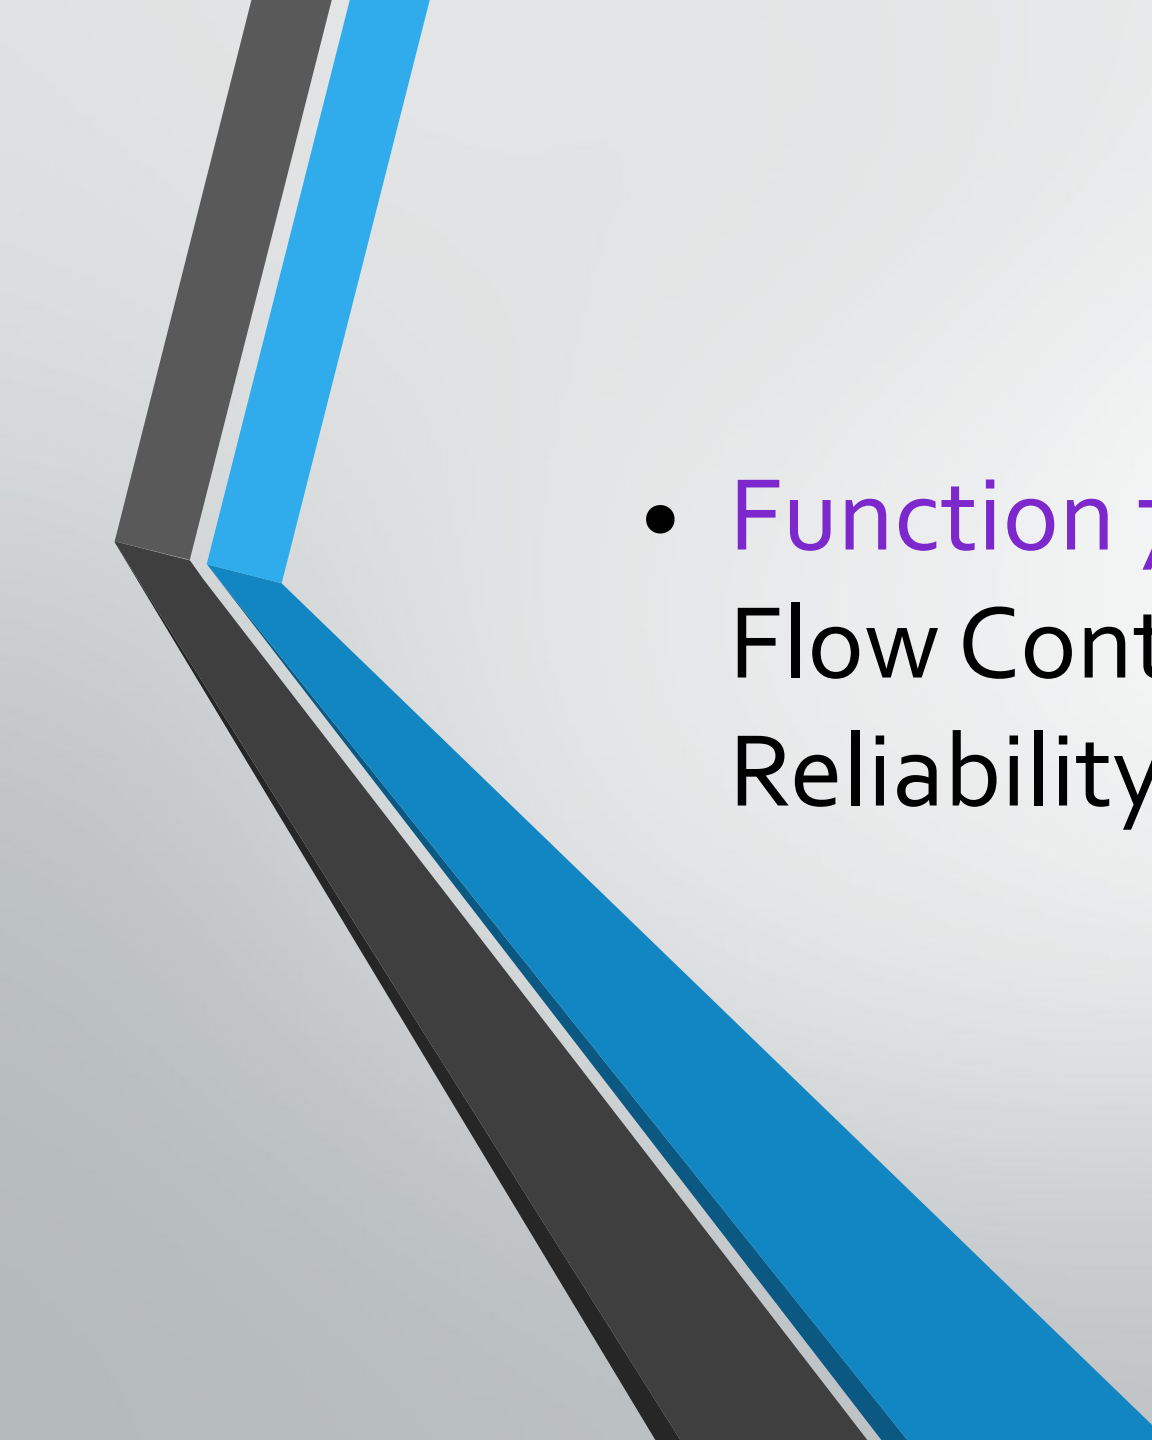
2

3

3 ACKs – Fast Transmission

Time Out Timer



- 
- **Function 7 :**
Flow Control and Recovery for
Reliability

Flow Control

- Transmission Control Protocol (TCP) uses a *sliding window* for flow control.
- What is the “**Window**”?
 - Indicates the size of the device's receive buffer for the particular connection.
 - How much data a device can handle from its peer at one time before it is passed to the application process.
 - Set by receiver of data
 - **Example** : The server's window size was **360**. This means the receiver is willing to take **no more than 360 bytes** at a time from the sender.

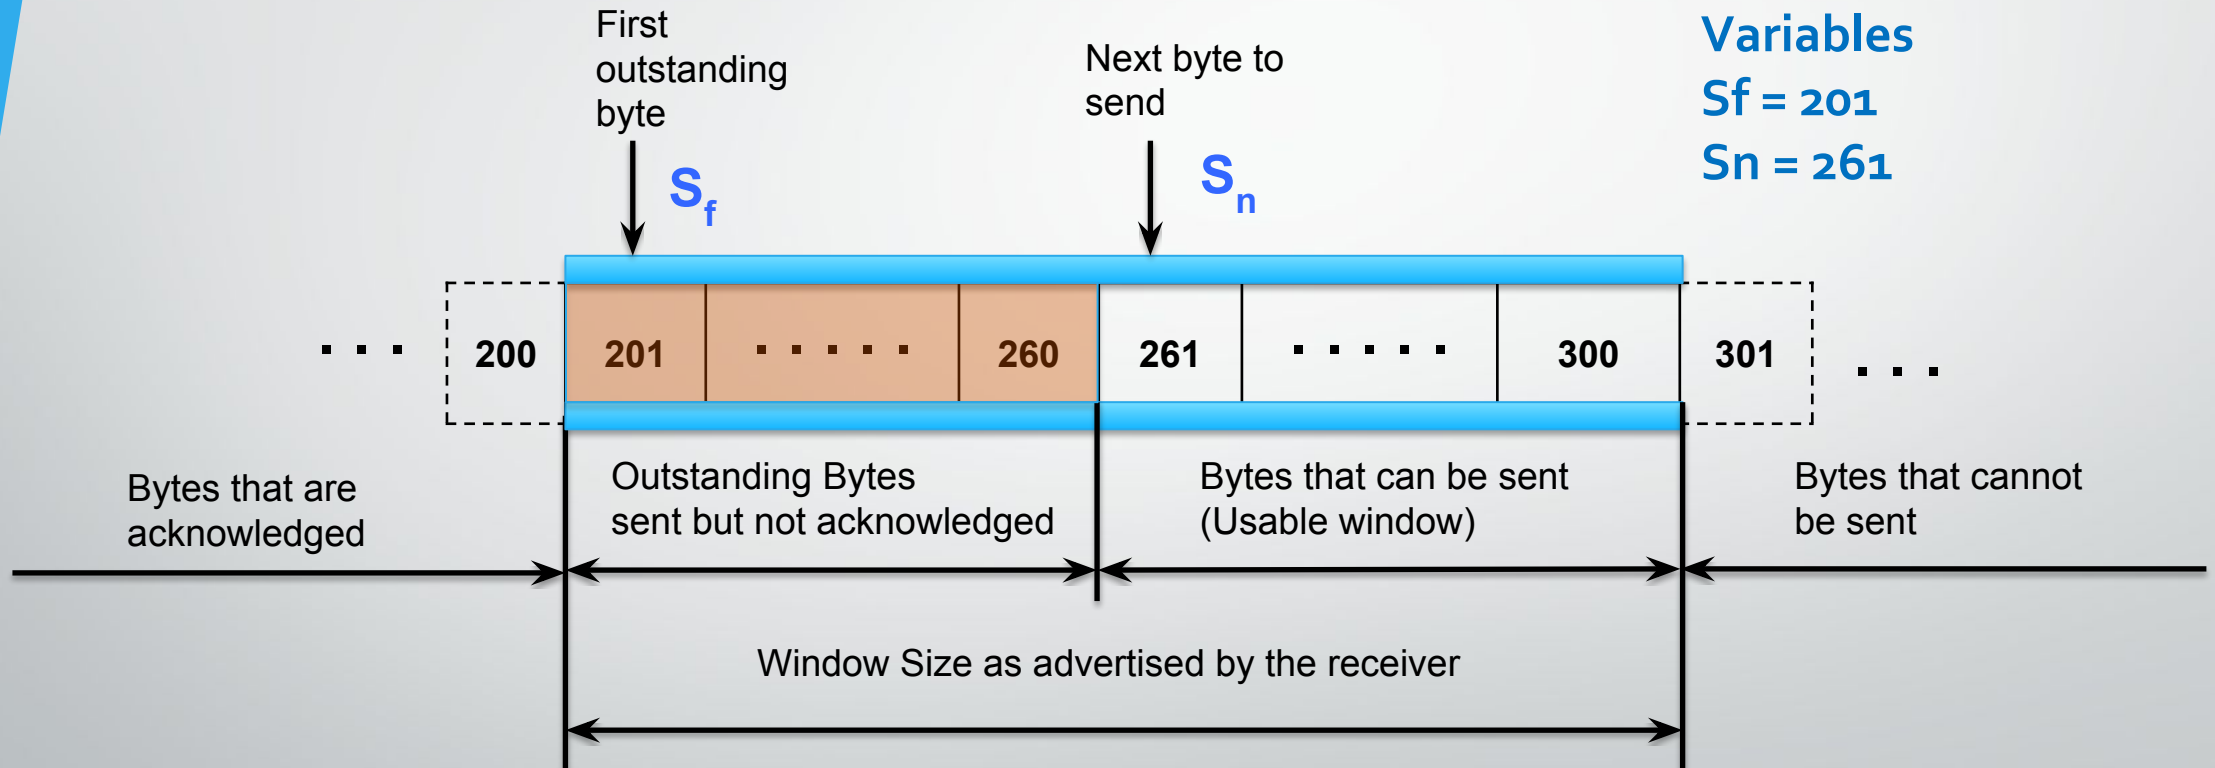
Sender Sliding Window

Window Size = 100 bytes

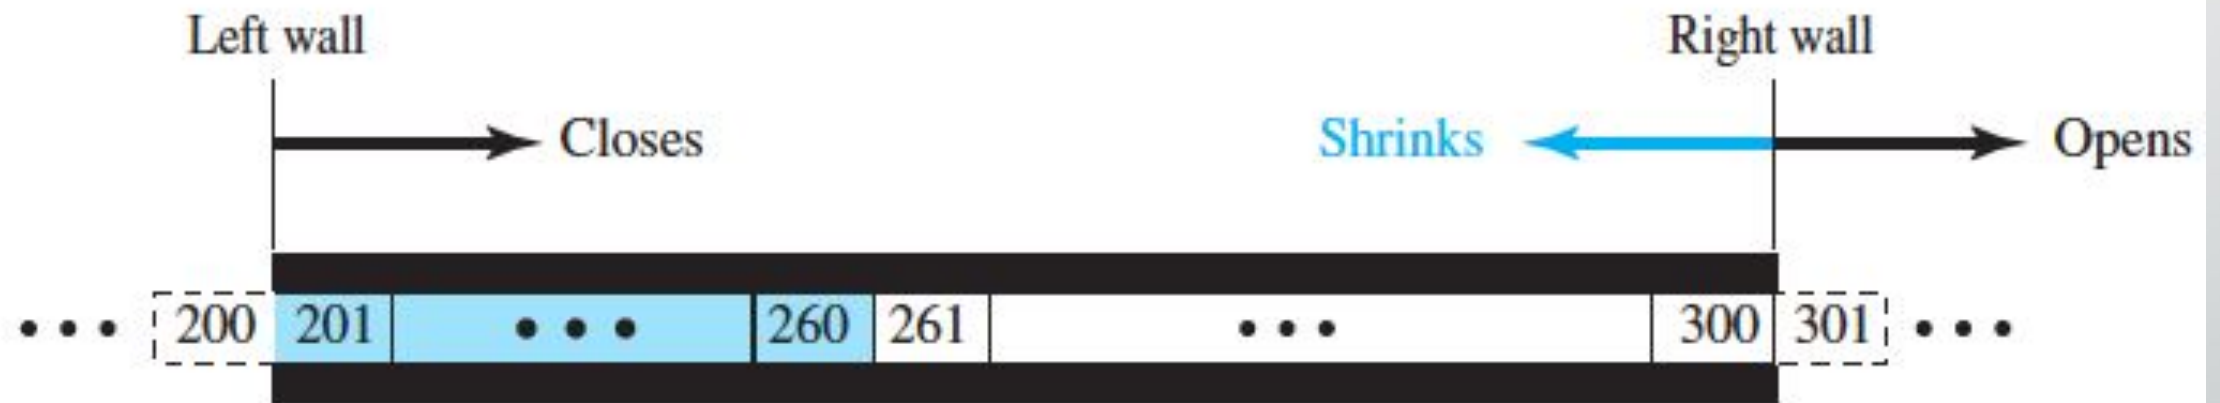
Variables

$S_f = 201$

$S_n = 261$



Sliding of Sender Window



b. Opening, closing, and shrinking send window

Sliding of Sender Window

#Sender receives a segment with ACK 241

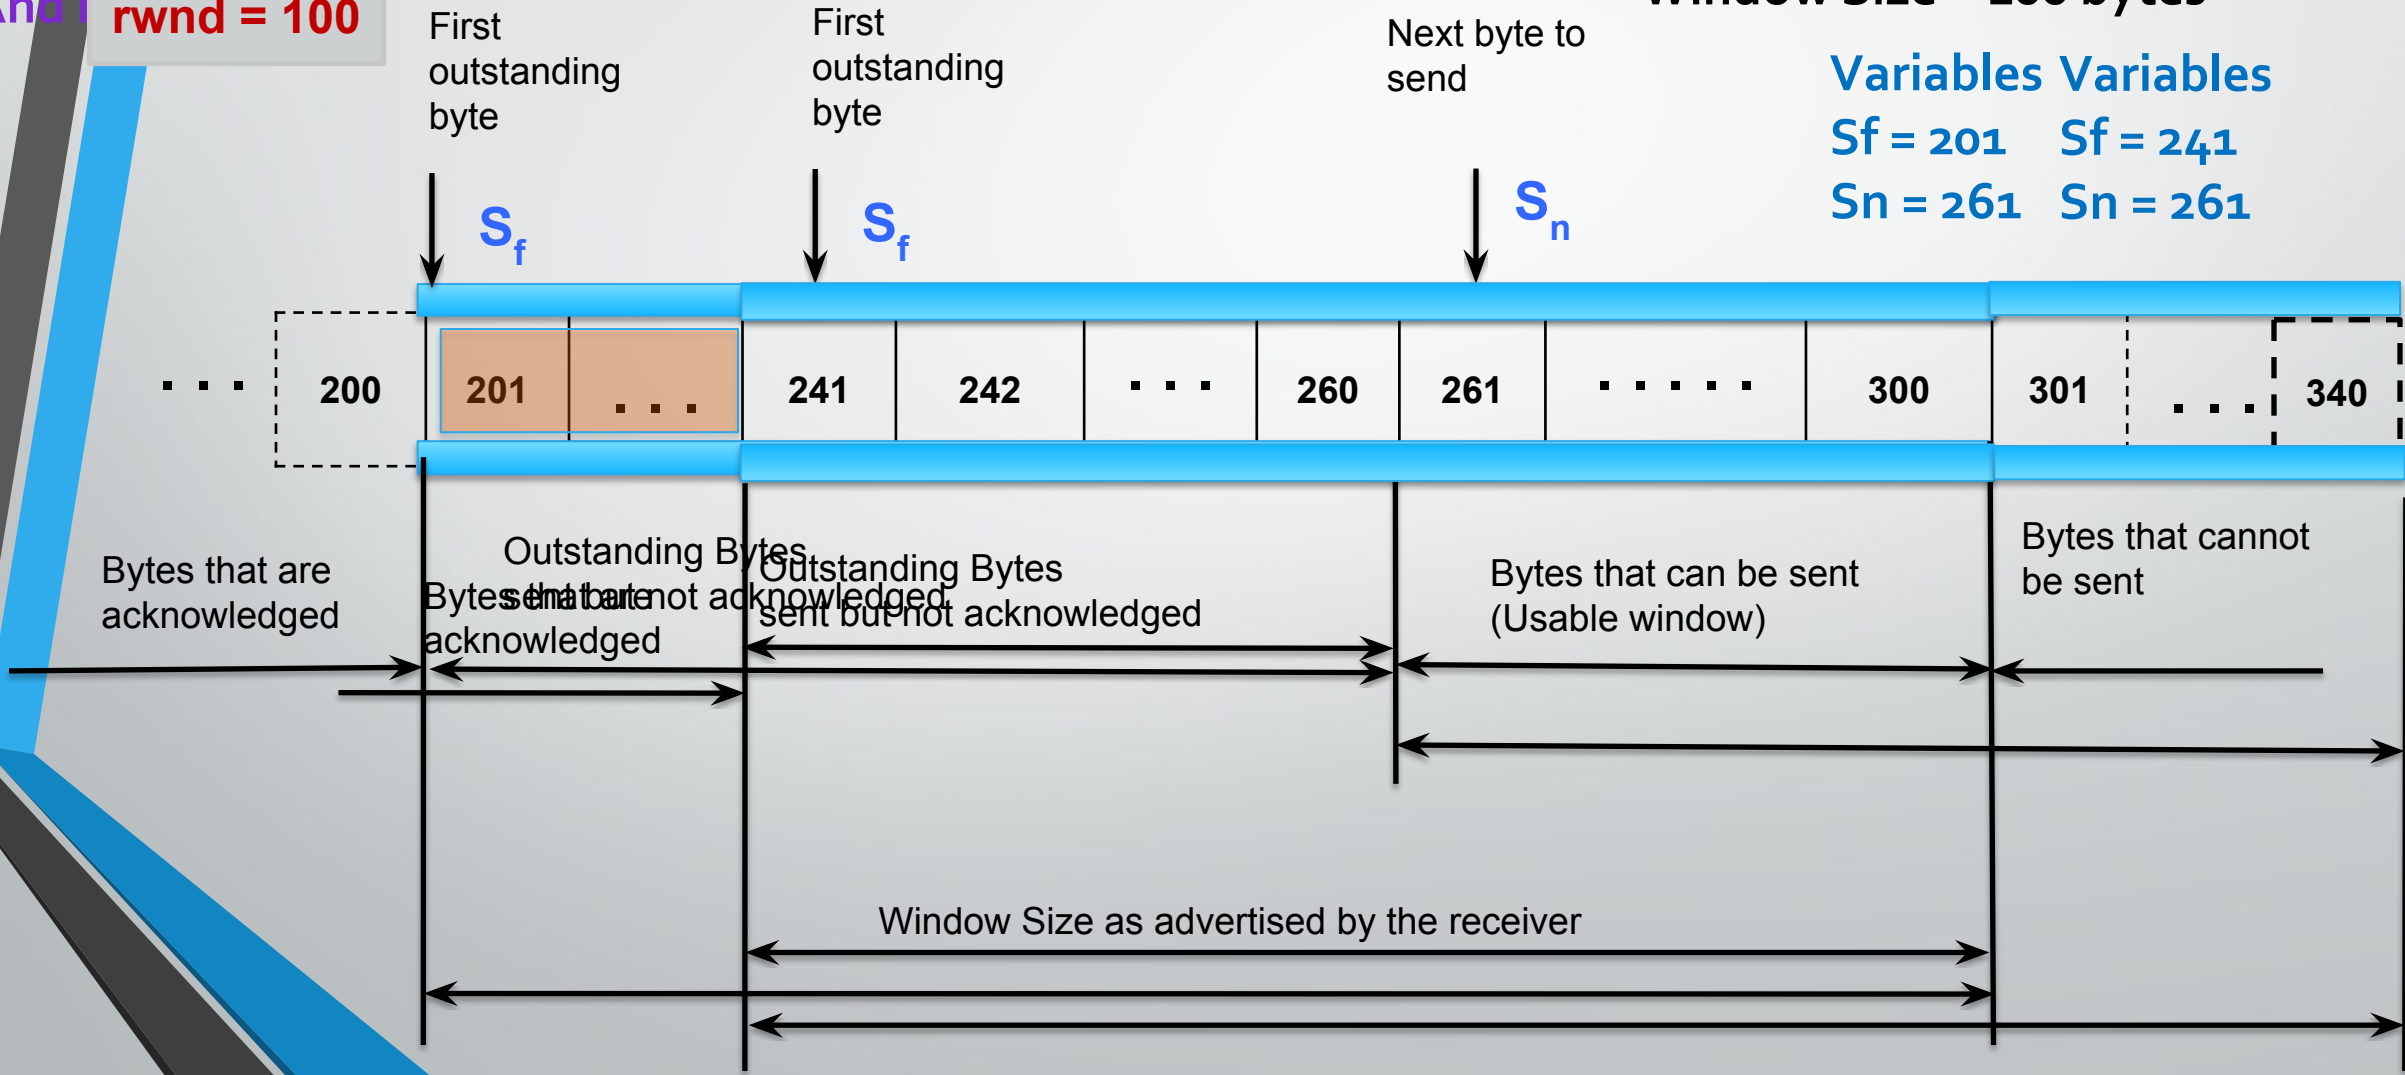
And $rwnd = 100$

Window Size = 100 bytes

Variables Variables

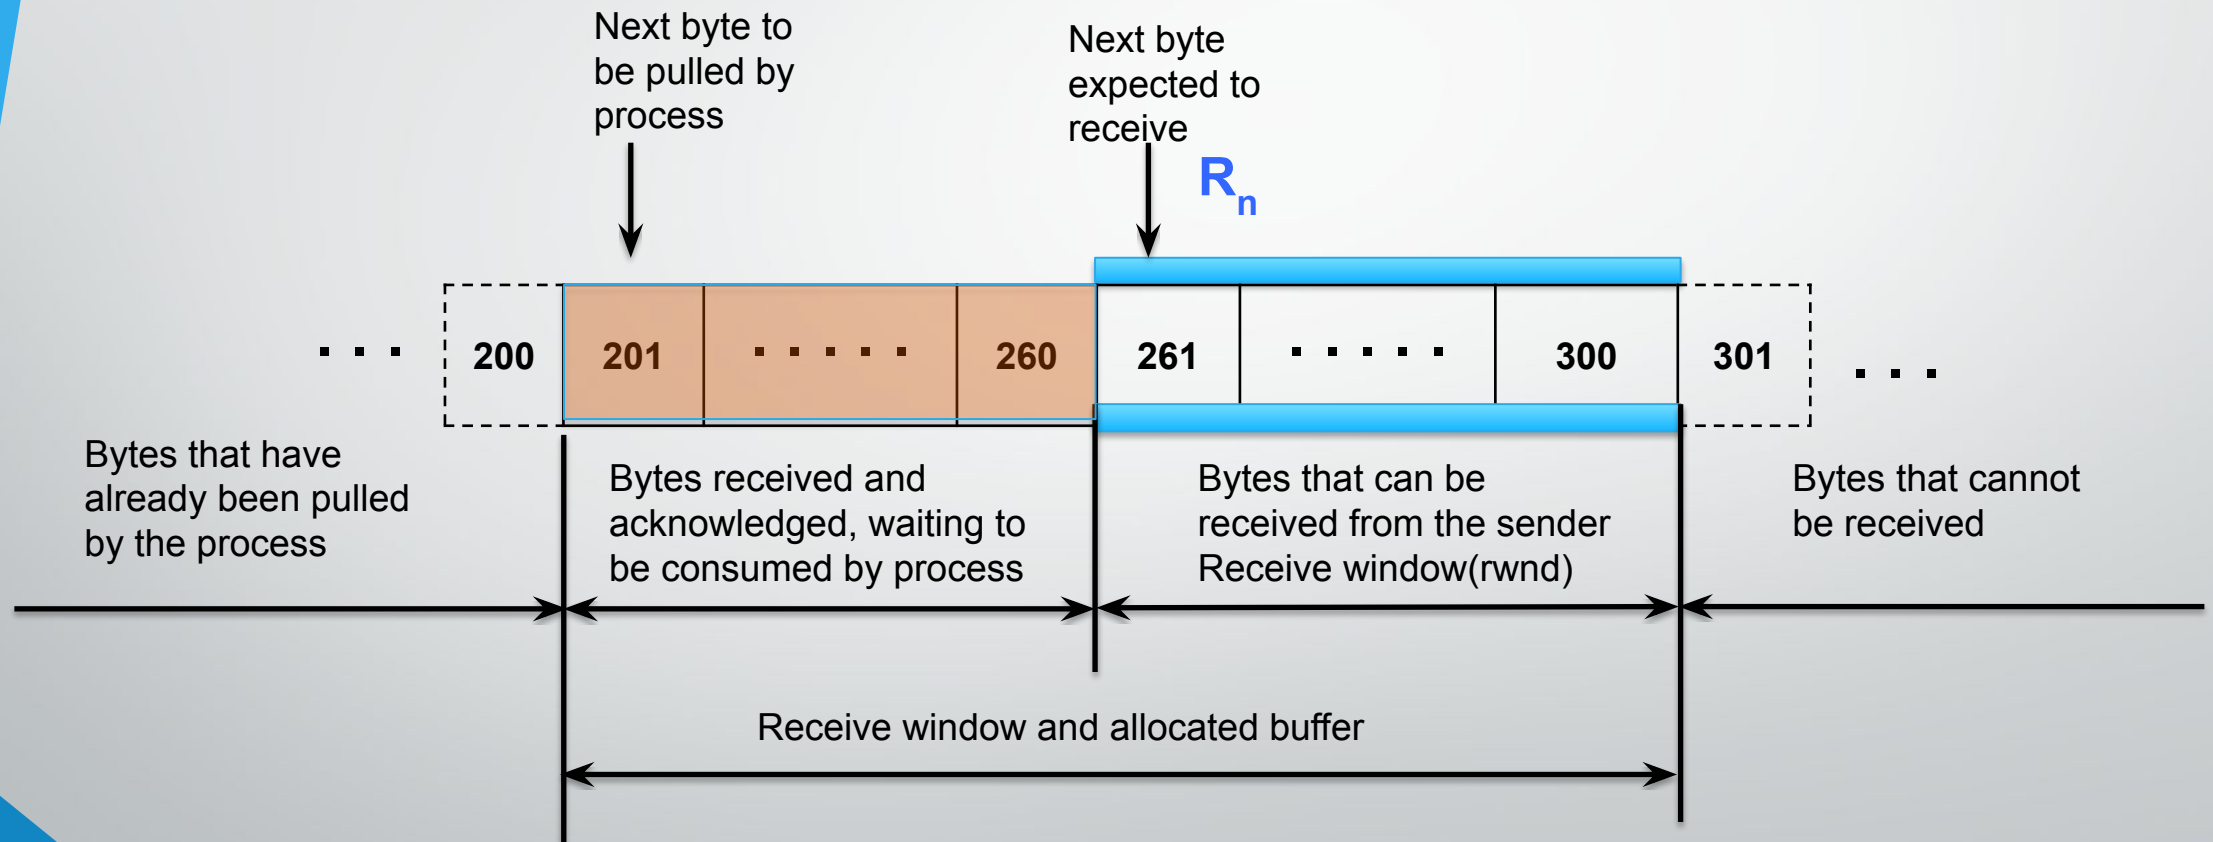
$Sf = 201$ $Sf = 241$

$Sn = 261$ $Sn = 261$




Receiver Sliding Window

Window Size = 100 bytes

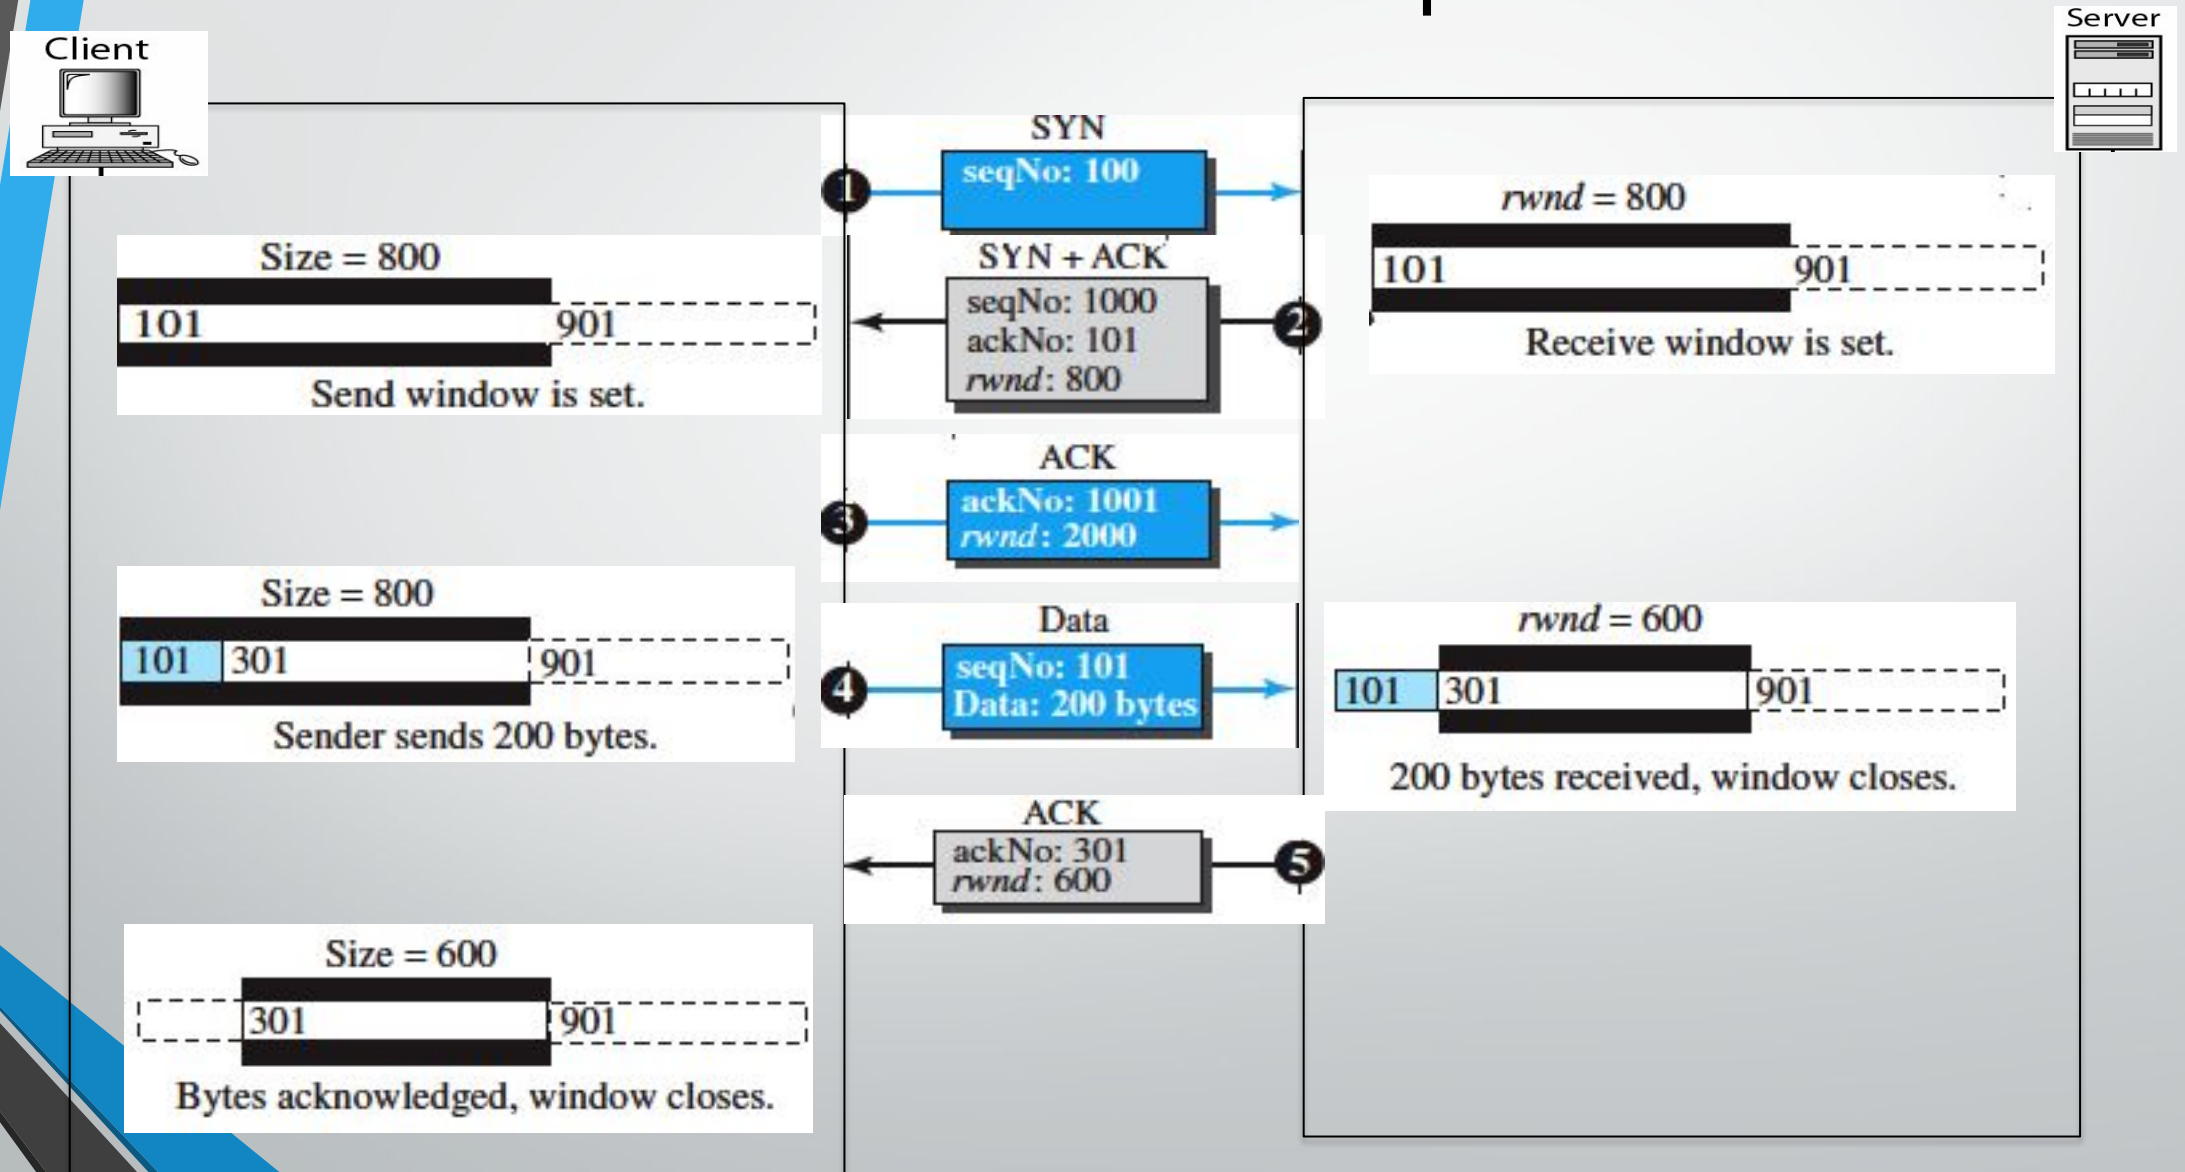



- $rwnd = \text{buffer size} - \text{number of bytes to be pulled} = 40 \text{ bytes}$

Sliding of Receiver Window



Flow Control Example






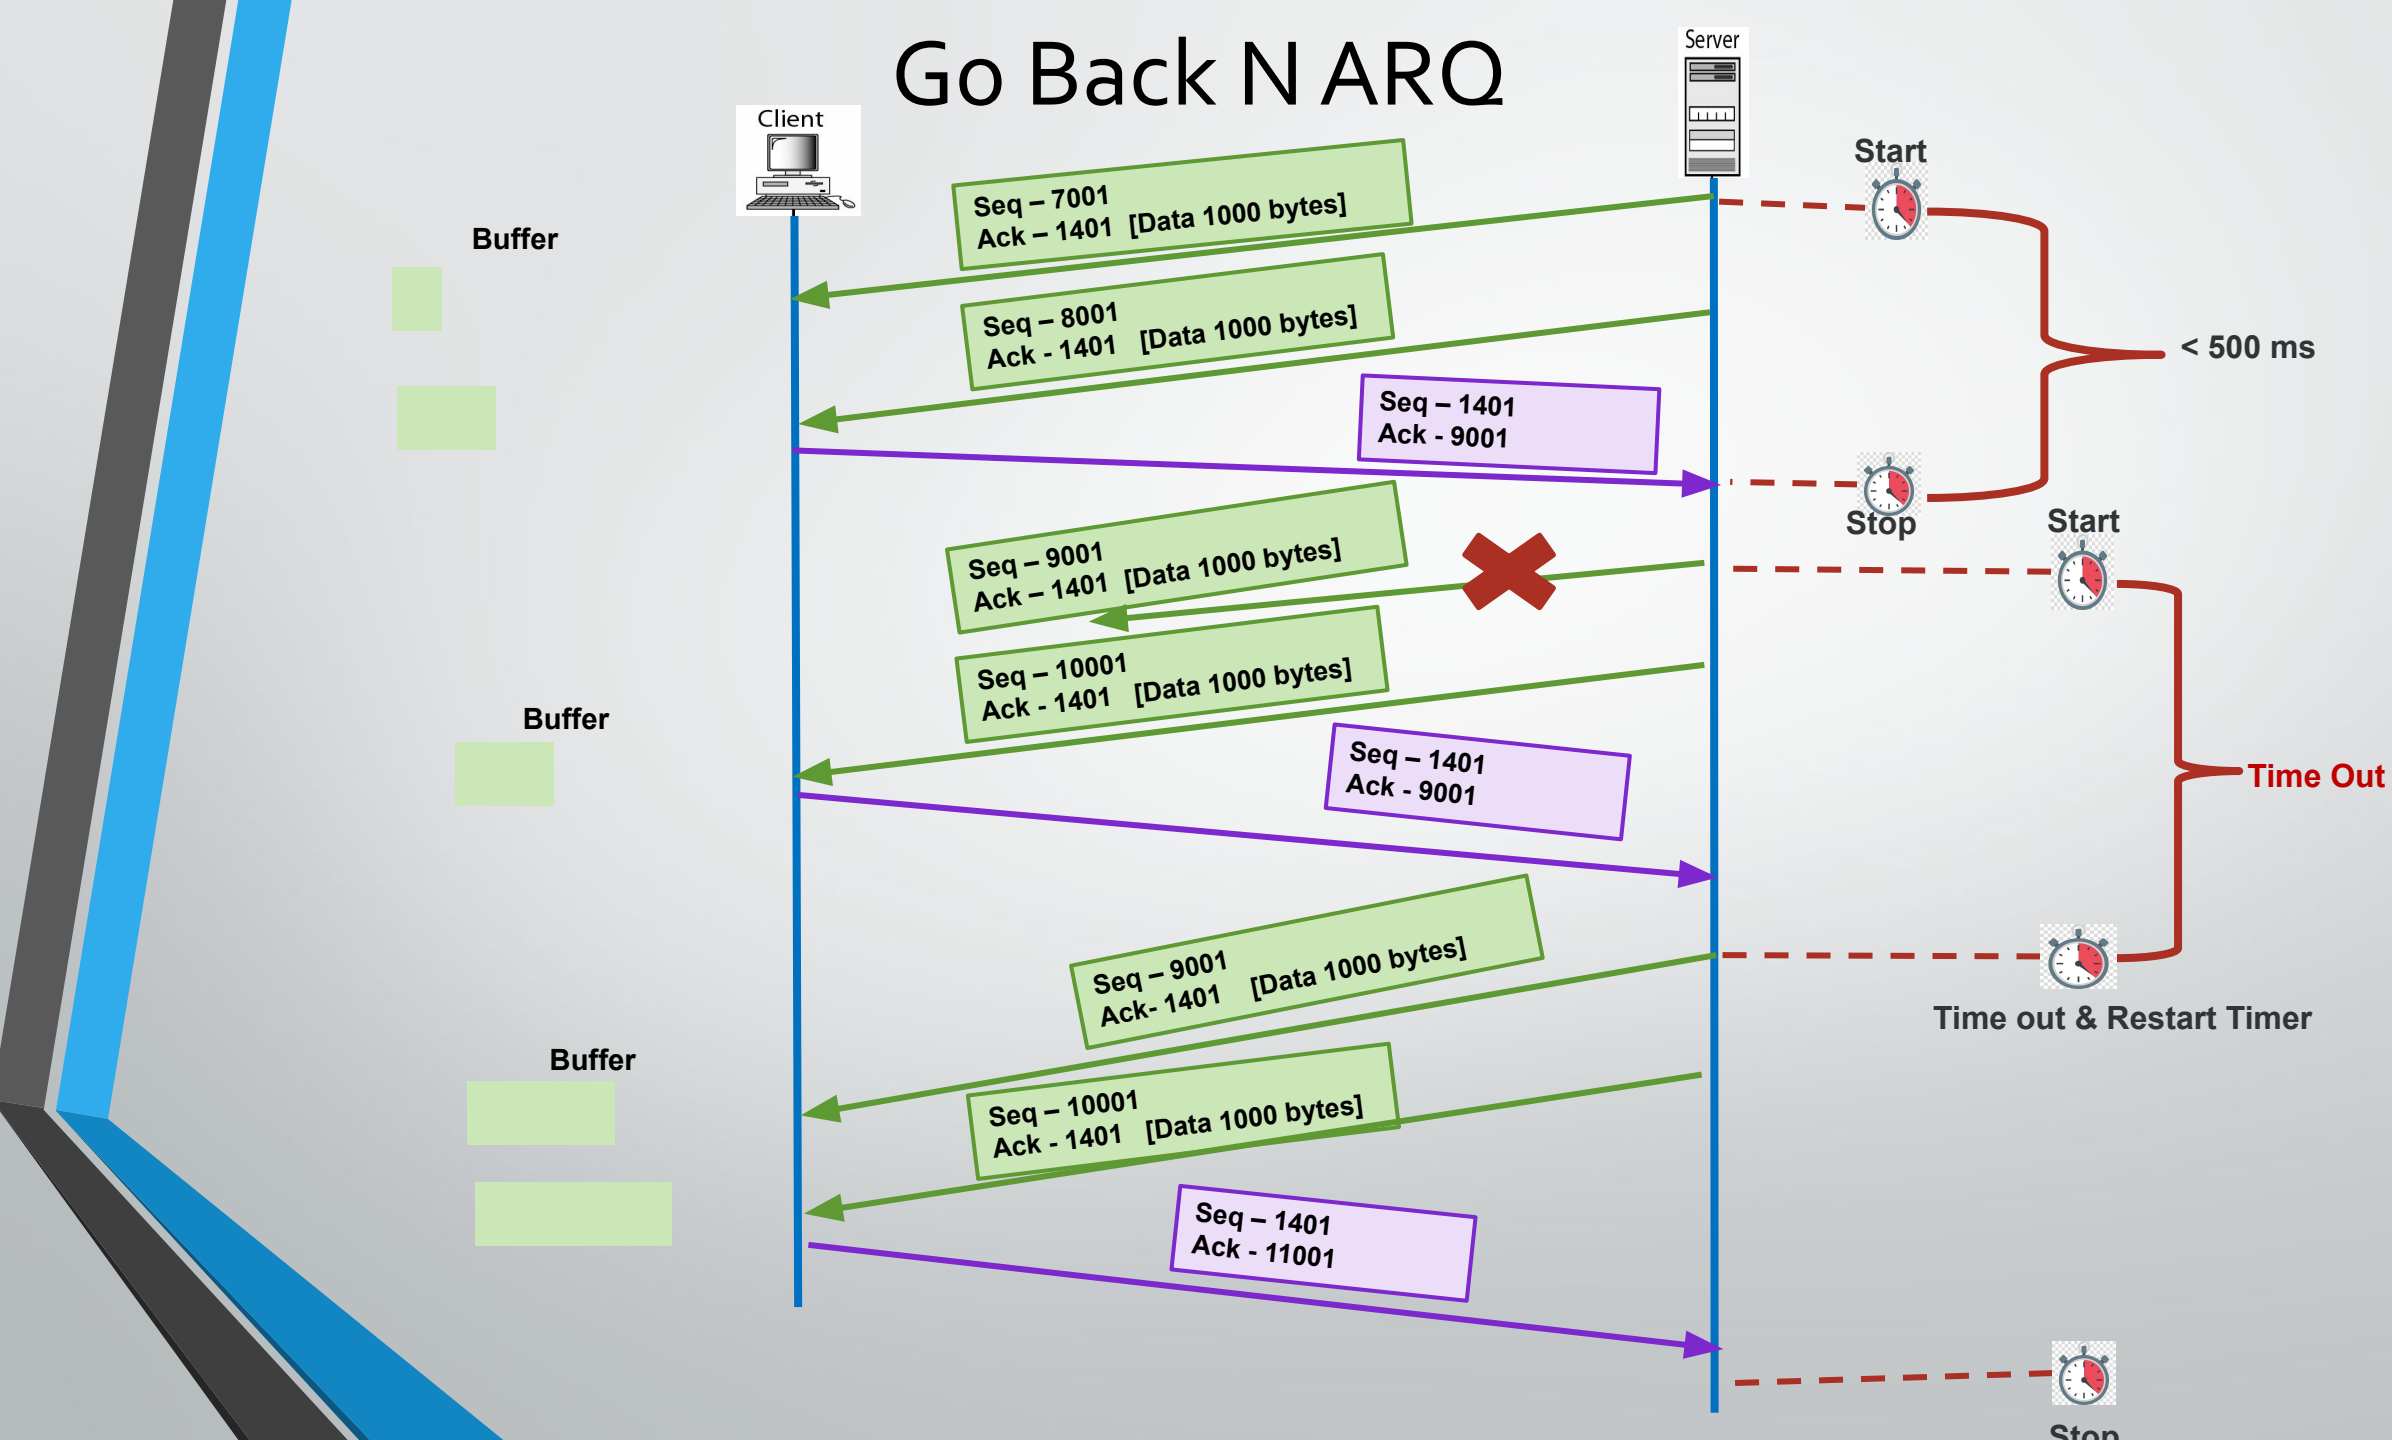
Different TCP Sliding Window Protocols

- **Selective Repeat Protocol**

- Only those segments are re-transmitted which are found lost or corrupted
- Keep track of out of order segments at the receiver side
- More efficient for noisy channels
- Widely used in TCP



Go Back N ARQ



Different TCP Sliding Window Protocols

- **Go Back N Protocol**

- If the sent segment are are found corrupted or lost then all the segments are re-transmitted from the lost segment to the last segment transmitted
- Do not keep track of out of order segments
- Efficient for less noisy channel

Overall Flow control

- The initial window size is agreed during the **three-way handshake**.
- If this is too much for the receiver and it **loses data** (e.g. buffer overflow) then it can **decrease** the window size.
- If **all is well** then the receiver will **increase** the window size.



The End